

## Python Machine Learning Mini Course

<https://machinelearningmastery.com/python-machine-learning-mini-course/>  
(<https://machinelearningmastery.com/python-machine-learning-mini-course/>)

I'm using a mix of descriptions used in the course along with my own commentary in code comments and markdown blocks to learn and help display understanding

## Lesson 01

```
In [1]: # Lesson 01
import sys
import scipy
import numpy as np
import matplotlib as plt
import pandas as pd
import sklearn as sk
# giving the imports nicknames; removed the prints to keep it cleaner
```

## Lesson 02

```
In [2]: # Lesson 02
myarray = np.array([[1, 2, 3], [4, 5, 6]])
rownames = ['a', 'b']
colnames = ['one', 'two', 'three']
mydataframe = pd.DataFrame(myarray, index=rownames, columns=colnames)

print(mydataframe)
```

	one	two	three
a	1	2	3
b	4	5	6

## Lesson 03

```
In [3]: # Lesson 03
url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indian-s-diabetes.data.csv"
names = ['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class']
data = pd.read_csv(url, names=names)

print(data.shape)
```

(768, 9)

names: list of the columns being imported/read as a csv

9 columns, 768 rows

Below are other methods of importing a csv file

```
In [4]: # import csv
# with open('example.csv', 'r') as file:
#     csv_reader = csv.reader(file)

# import numpy as np
# data = np.loadtxt('example.csv', delimiter=',')

# import pandas as pd
# data = pd.read_csv('example.csv')
```

## Lesson 04

```
In [5]: # Lesson 04
# Statistical Summary
data = pd.read_csv(url, names=names)
data.describe()
```

Out[5]:

	preg	plas	pres	skin	test	mass	pedi	
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.234516
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.767350
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.010000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000

```
In [6]: data.head()
```

Out[6]:

	preg	plas	pres	skin	test	mass	pedi	age	class
<b>0</b>	6	148	72	35	0	33.6	0.627	50	1
<b>1</b>	1	85	66	29	0	26.6	0.351	31	0
<b>2</b>	8	183	64	0	0	23.3	0.672	32	1
<b>3</b>	1	89	66	23	94	28.1	0.167	21	0
<b>4</b>	0	137	40	35	168	43.1	2.288	33	1

In [7]: `data.shape`

Out[7]: (768, 9)

In [8]: `data.dtypes`

Out[8]:

```

preg      int64
plas      int64
pres      int64
skin      int64
test      int64
mass      float64
pedi      float64
age       int64
class     int64
dtype: object

```

In [9]: `data.describe()`

Out[9]:

	preg	plas	pres	skin	test	mass	pedi	age	class
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.233416	0.511368
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.768013	0.500879
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.010000	0.000000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	0.000000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	0.000000

In [10]: `data.corr()`


Out[10]:

	preg	plas	pres	skin	test	mass	pedi	age	class
<b>preg</b>	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0.544341	0.221898
<b>plas</b>	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514	0.466581
<b>pres</b>	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528	0.065068
<b>skin</b>	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970	0.074752
<b>test</b>	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163	0.130548
<b>mass</b>	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0.036242	0.292695
<b>pedi</b>	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0.033561	0.173844
<b>age</b>	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.033561	1.000000	0.238356
<b>class</b>	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.173844	0.238356	1.000000

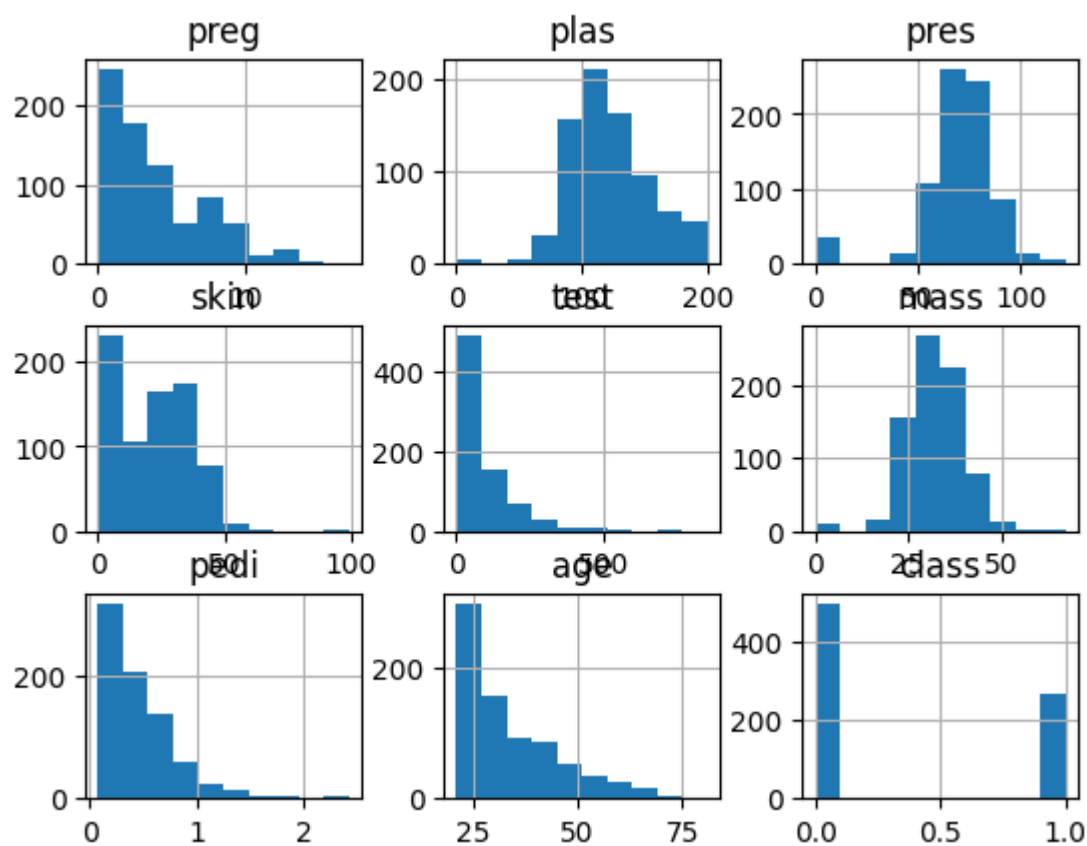
## Lesson 05

```
In [11]: # Lesson 05
# Scatter Plot Matrix
import matplotlib.pyplot as plt
from pandas.plotting import scatter_matrix


scatter_matrix(data)
plt.show()
```



```
In [12]: # Histogram of the data  
data.hist()  
plt.show()
```



```
In [13]: # Box plot of the data  
data.plot(kind='box')  
plt.show()
```



## Lesson 06

```
In [14]: # Lesson 06
# Prepare the data for modeling by preprocessing it
# Standardize data (0 mean, 1 stdev)
from sklearn.preprocessing import StandardScaler

dataframe = pd.read_csv(url, names=names)
array = dataframe.values

# separate array into input and output components
X = array[:,0:8]
Y = array[:,8]

# calculate params needed to standardize data
scaler = StandardScaler().fit(X)
rescaledX = scaler.transform(X)

# summarize transformed data
np.set_printoptions(precision=3)

print(rescaledX[0:5,:])

[[ 0.64  0.848  0.15  0.907 -0.693  0.204  0.468  1.426]
 [-0.845 -1.123 -0.161  0.531 -0.693 -0.684 -0.365 -0.191]
 [ 1.234  1.944 -0.264 -1.288 -0.693 -1.103  0.604 -0.106]
 [-0.845 -0.998 -0.161  0.155  0.123 -0.494 -0.921 -1.042]
 [-1.142  0.504 -1.505  0.907  0.766  1.41  5.485 -0.02 ]]
```

The above snippet:

- Loads the Pima Indians onset of diabetes dataset
- Calculates the parameters needed to standardize the data, transforming the data to have a mean of 0 and a standard deviation of 1
- Then creates a standardized copy of the input data. `StandardScaler()` is used to do that and assign it to a new variable

```
In [15]: # Standardize numerical data (e.g. mean of 0 and standard deviation of 1) using the scale and center options.

from sklearn.preprocessing import MinMaxScaler

# Normalize data (0 to 1)
minmax_scaler = MinMaxScaler(feature_range=(0, 1))
normalizedX = minmax_scaler.fit_transform(X)

# Summarize transformed data
print(normalizedX[0:5,:])

[[0.353 0.744 0.59  0.354 0.    0.501 0.234 0.483]
 [0.059 0.427 0.541 0.293 0.    0.396 0.117 0.167]
 [0.471 0.92  0.525 0.    0.    0.347 0.254 0.183]
 [0.059 0.447 0.541 0.232 0.111 0.419 0.038 0.    ]
 [0.    0.688 0.328 0.354 0.199 0.642 0.944 0.2   ]]
```

The difference between `MinMaxScaler()` and `StandardScaler()` is that `MinMaxScaler` will standardize the data between a range of [0,1] while `StandardScaler` will standardize it so it has a mean of 0 and standard deviation of 1

## Lesson 07

```
In [16]: # Lesson 07

# Evaluate using Cross Validation
from pandas import read_csv
from sklearn.model_selection import KFold, cross_val_score
from sklearn.linear_model import LogisticRegression

dataframe = read_csv(url, names=names)
array = dataframe.values
# split the data's features and target values
X = array[:,0:8]
Y = array[:,8]
kfold = KFold(n_splits=10, random_state=7, shuffle=True)
model = LogisticRegression(solver='liblinear')
# test and training sets are handled by cross_val_score
results = cross_val_score(model, X, Y, cv=kfold)

print(f"Accuracy: {results.mean()*100:.3f}% ({results.std()*100:.3f}%)")

Accuracy: 77.086% (5.091%)
```



The lesson suggests to attempt the following actions:

- Split a dataset into training and test sets.
- Estimate the accuracy of an algorithm using k-fold cross validation.
- Estimate the accuracy of an algorithm using leave one out cross validation.

It's kind of there already, just a few minor changes to be made

```
In [17]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_m
atrix

# Split the data into training and test sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.20, rand
om_state=7)

# Train the model on the training set
model = LogisticRegression(solver='liblinear')
model.fit(X_train, Y_train)

# Predict the outcomes on the test set
Y_pred = model.predict(X_test)

# Evaluate the model's performance
accuracy = accuracy_score(Y_test, Y_pred)
print(f"Accuracy: {accuracy * 100:.3f}%")
print(f"\nClassification Report:\n{classification_report(Y_test, Y_pred)}")
print(f"\nConfusion Matrix:\n{confusion_matrix(Y_test, Y_pred)}")
```

Accuracy: 79.221%

Classification Report:

	precision	recall	f1-score	support
0.0	0.78	0.94	0.85	97
1.0	0.84	0.54	0.66	57
accuracy			0.79	154
macro avg	0.81	0.74	0.76	154
weighted avg	0.80	0.79	0.78	154

Confusion Matrix:

```
[[91  6]
 [26 31]]
```

A similar method for the above cell was used in Step 5.1 of the other assignment. I also came back to this after Lesson 08 to apply the evaluation metrics

[https://scikit-learn.org/stable/modules/generated/sklearn.metrics.accuracy\\_score.html](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.accuracy_score.html) ([https://scikit-learn.org/stable/modules/generated/sklearn.metrics.accuracy\\_score.html](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.accuracy_score.html))

In [18]: *# Estimate the accuracy of an algorithm using k-fold cross validation*

```
kfold = KFold(n_splits=10, random_state=7, shuffle=True)
results = cross_val_score(model, X, Y, cv=kfold)
print(results)
```

[0.831 0.714 0.714 0.779 0.792 0.766 0.688 0.857 0.803 0.763]

In [19]: **from sklearn.model\_selection import** LeaveOneOut

```
# Estimate the accuracy of an algorithm using Leave one out cross validation.
loo = LeaveOneOut()
results_loo = cross_val_score(model, X, Y, scoring='accuracy', cv=loo)
print(f"Accuracy: {results_loo.mean()*100:.3f}% ({results_loo.std()*100:.3f}%)"
```

Accuracy: 76.823% (42.196%)

[https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.LeaveOneOut.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.LeaveOneOut.html) ([https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.LeaveOneOut.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.LeaveOneOut.html))

Example under "LOOCV to Evaluate Machine Learning Models" section

<https://machinelearningmastery.com/loocv-for-evaluating-machine-learning-algorithms/>  
(<https://machinelearningmastery.com/loocv-for-evaluating-machine-learning-algorithms/>)

## Lesson 08

In [20]: *# Lesson 08*

```
# Cross Validation Classification LogLoss
dataframe = read_csv(url, names=names)
array = dataframe.values
X = array[:,0:8]
Y = array[:,8]
# I added shuffle=True to the kFold, or else it would error out
kfold = KFold(n_splits=10, random_state=7, shuffle=True)
model = LogisticRegression(solver='liblinear')
scoring = 'neg_log_loss'
results = cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
# I had to move one of the ')' to the end. Originally it was in the wrong location
print(f"Logloss: {results.mean():.3f} ({results.std():.3f})")
```

Logloss: -0.494 (0.042)

- Practice using the Accuracy and LogLoss metrics on a classification problem.
- Practice generating a confusion matrix and a classification report.
- Practice using RMSE and RSquared metrics on a regression problem.

```
In [21]: # Practice using the Accuracy and LogLoss metrics on a classification problem.
results_acc = cross_val_score(model, X, Y, cv=kfold, scoring='accuracy')
print(f"Logloss: {results.mean()*100:.3f}% ({results.std()*100:.3f}%)" )
```

Logloss: -49.367% (4.207%)

[https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.cross\\_val\\_score.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.cross_val_score.html) ([https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.cross\\_val\\_score.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.cross_val_score.html))

[https://scikit-learn.org/stable/modules/model\\_evaluation.html](https://scikit-learn.org/stable/modules/model_evaluation.html) ([https://scikit-learn.org/stable/modules/model\\_evaluation.html](https://scikit-learn.org/stable/modules/model_evaluation.html))

Above is a link to different options accepted in scoring for model evaluation.

```
In [22]: from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split

# Practice generating a confusion matrix and a classification report.
# This was done in the last cell of the other assignment

# X and Y variables are created a couple cells above
# create an 80/20 split training and test
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.20, random_state=7)
model.fit(X_train, Y_train)
predictions = model.predict(X_test)

print(f"\nClassification Report:\n{classification_report(Y_test, predictions)}")
print(f"\nConfusion Matrix:\n{confusion_matrix(Y_test, predictions)}")
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.78	0.94	0.85	97
1.0	0.84	0.54	0.66	57
accuracy			0.79	154
macro avg	0.81	0.74	0.76	154
weighted avg	0.80	0.79	0.78	154

Confusion Matrix:

```
[[91  6]
 [26 31]]
```

```
In [23]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# practice using RMSE and RSquared metrics on a regression problem.

# references:
# https://machinelearningmastery.com/regression-metrics-for-machine-learning/
# https://www.statology.org/r-squared-in-python/

regression_model = LinearRegression()

# fit the regression model
regression_model.fit(X_train, Y_train)
Y_pred = regression_model.predict(X_test)

# calculate RMSE
rmse = mean_squared_error(Y_test, Y_pred, squared=False)

# calculate RSquared of regression model
r_squared = regression_model.score(X_train, Y_train)

print(f"RSquared: \t{r_squared:.3f}\nRMSE: \t\t{rmse:.3f}")
```

```
RSquared:      0.298
RMSE:          0.401
```

## Lesson 09

```
In [24]: # Lesson 09

# KNN Regression
from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsRegressor
url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/housing.data"
names = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV']
dataframe = read_csv(url, delim_whitespace=True, names=names)
array = dataframe.values
# split into features and target variables
X = array[:,0:13]
Y = array[:,13]
# shuffle=True was required to be added again
kfold = KFold(n_splits=10, random_state=7, shuffle=True)
model = KNeighborsRegressor()
# note: i linked list of accepted scoring strings in Lesson 08 along with what each does
# Mean squared error regression loss
scoring = 'neg_mean_squared_error'
results = cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
print(results.mean())

-38.852320266666666
```

```
In [38]: from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

# Spot check linear algorithms on a dataset (e.g. linear regression, logistic regression and linear discriminate analysis).

checks = {
    LinearRegression(): 'neg_mean_squared_error',
    LogisticRegression(solver='liblinear'): None,
    LinearDiscriminantAnalysis(): 'accuracy'
}
# Loop through above to apply the estimator and scoring to the X and Y created earlier
for check, score in checks.items():
    results = cross_val_score(check, X, Y, cv=kfold, scoring=score)
    print(f"{str(check).split('(')[0]}: {results.mean()}")

LinearRegression: -0.1632073857590707
LogisticRegression: 0.7708646616541353
LinearDiscriminantAnalysis: 0.7669685577580315
```

```
In [26]: from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor

# Spot check some non-Linear algorithms on a dataset (e.g. KNN, SVM and CART).

checks = {
    SVR(): 'neg_mean_squared_error',
    DecisionTreeRegressor(): 'neg_mean_squared_error',
    KNeighborsRegressor(): 'neg_mean_squared_error'
}
for check, score in checks.items():
    results = cross_val_score(check, X, Y, cv=kfold, scoring=score)
    print(f"{str(check).split('(')[0]}: {results.mean()}")
```

SVR: -67.64140705473743

DecisionTreeRegressor: -20.25612431372549

KNeighborsRegressor: -38.852320266666666

For spot checking algorithms, I don't know if using the same scoring is required throughout. Using certain combinations of algo and evaluation score would cause errors. Atleast with the method of application I am using, it is easy to update with the correct value.

## Lesson 10

```
In [27]: # Lesson 10

# Compare Algorithms
from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
# Load dataset
url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indian-s-diabetes.data.csv"
names = ['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class']
dataframe = read_csv(url, names=names)
array = dataframe.values
X = array[:,0:8]
Y = array[:,8]
# prepare models to be compared
linear_models = [('LR', LogisticRegression(solver='liblinear')),
                  ('LDA', LinearDiscriminantAnalysis())]
# evaluate each model in turn
results = []
names = []
scoring = 'accuracy'
for name, model in linear_models:
    kfold = KFold(n_splits=10, random_state=7, shuffle=True)
    cv_results = cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    # i hate original way of printing vars compared to using f-strings
    msg = f"{name}: {cv_results.mean():.6f} ({cv_results.std():.6f})"
    print(msg)
```

LR: 0.770865 (0.050905)

LDA: 0.766969 (0.047966)

- Compare linear algorithms to each other on a dataset.
- Compare nonlinear algorithms to each other on a dataset.
- Compare different configurations of the same algorithm to each other.
- Create plots of the results comparing algorithms.

## Linear

```

LinearRegression(): 'neg_mean_squared_error'
LogisticRegression(solver='liblinear'): None
LinearDiscriminantAnalysis(): 'accuracy'

```

## Non Linear

```

SVR(): 'neg_mean_squared_error'
DecisionTreeRegressor(): 'neg_mean_squared_error'
KNeighborsRegressor(): 'neg_mean_squared_error'

```

## Different configurations

```

LinearRegression(): 'neg_mean_squared_error'
LinearRegression(): 'accuracy'

```

```

In [28]: nonlinear_models = [('SVR', SVR()),
                             ('DTR', DecisionTreeRegressor())]

results = []
names = []
scoring = 'neg_mean_squared_error'
for name, model in nonlinear_models:
    kfold = KFold(n_splits=10, random_state=7, shuffle=True)
    cv_results = cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    # i hate original way of printing vars compared to using f-strings
    msg = f"{name}: {cv_results.mean():.6f} ({cv_results.std():.6f})"
    print(msg)

```

SVR: -0.174617 (0.020384)

DTR: -0.305878 (0.048066)



```
In [29]: config_model = [('LR newton-cg', LogisticRegression(solver='newton-cg')),
                        ('LR liblinear', LogisticRegression(solver='liblinear'))]
results = []
names = []
scoring = 'neg_mean_squared_error'
# Loop through each option, both using LogisticRegression but with different solvers
for name, model in config_model:
    kfold = KFold(n_splits=10, random_state=7, shuffle=True)
    cv_results = cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    # i hate original way of printing vars compared to using f-strings
    msg = f"{name}: {cv_results.mean():.6f} ({cv_results.std():.6f})"
    print(msg)
```

LR newton-cg: -0.227837 (0.049684)

LR liblinear: -0.229135 (0.050905)

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LogisticRegression.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html) ([https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LogisticRegression.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html))

## Lesson 11

```
In [30]: # Lesson 11

# Grid Search for Algorithm Tuning
from pandas import read_csv
import numpy
from sklearn.linear_model import Ridge
from sklearn.model_selection import GridSearchCV

# Load the dataset
url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indian-s-diabetes.data.csv"
names = ['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class']
dataframe = read_csv(url, names=names)
array = dataframe.values
X = array[:,0:8]
Y = array[:,8]

# define alpha values for the grid
alphas = numpy.array([1,0.1,0.01,0.001,0.0001,0])
param_grid = dict(alpha=alphas)

# grid search for ridge regression algo
model = Ridge()
grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=3)
grid.fit(X, Y)
print(grid.best_score_)
print(grid.best_estimator_.alpha)

0.27961755931297233
1.0
```

- Tune the parameters of an algorithm using a grid search that you specify.
- Tune the parameters of an algorithm using a random search.

```
In [31]: from sklearn.model_selection import RandomizedSearchCV

# define alpha values for the grid
alphas = numpy.array([1,0.1,0.01,0.001,0.0001,0])
param_random = dict(alpha=alphas)

# grid search for ridge regression algo
model = Ridge()

# setup random search with 3-fold cross validation
random_search = RandomizedSearchCV(estimator=model, param_distributions=param_
random, n_iter=6)
random_search.fit(X, Y)
# not defining n_iter gives a warning because it defaults to 10 while there is
6 alpha values

print(random_search.best_score_)
print(random_search.best_estimator_.alpha)

0.27610844129292433
1.0
```

[https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.RandomizedSearchCV.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.RandomizedSearchCV.html)  
[\(https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.RandomizedSearchCV.html\)](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.RandomizedSearchCV.html)

## Lesson 12

```
In [32]: # Lesson 12

# Random Forest Classification
from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indian
s-diabetes.data.csv"
names = ['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'clas
s']
dataframe = read_csv(url, names=names)
array = dataframe.values
X = array[:,0:8]
Y = array[:,8]
num_trees = 100
max_features = 3
kfold = KFold(n_splits=10, random_state=7, shuffle=True)
model = RandomForestClassifier(n_estimators=num_trees, max_features=max_featur
es)
results = cross_val_score(model, X, Y, cv=kfold)
print(results.mean())

0.7643369788106631
```

- Practice bagging ensembles with the random forest and extra trees algorithms.
- Practice boosting ensembles with the gradient boosting machine and AdaBoost algorithms.
- Practice voting ensembles using by combining the predictions from multiple models together.

```
In [33]: from sklearn.ensemble import ExtraTreesClassifier
model_extratrees = ExtraTreesClassifier(n_estimators=num_trees, max_features=m
ax_features)
results_extratrees = cross_val_score(model_extratrees, X, Y, cv=kfold)
print(f"Extra Trees: {results_extratrees.mean()}")
```

Extra Trees: 0.7682330827067669

```
In [34]: from sklearn.ensemble import GradientBoostingClassifier
model_gradient = GradientBoostingClassifier(n_estimators=num_trees, max_featur
es=max_features)
results_gradient = cross_val_score(model_gradient, X, Y, cv=kfold)
print(f"Gradient Boosting: {results_gradient.mean()}")
```

Gradient Boosting: 0.764354066985646

```
In [35]: from sklearn.ensemble import AdaBoostClassifier
model_ada = AdaBoostClassifier(n_estimators=num_trees)
results_ada = cross_val_score(model_ada, X, Y, cv=kfold)
print(f"AdaBoost: {results_ada.mean()}")
```

AdaBoost: 0.7578605604921395

Extra Trees: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.ExtraTreesClassifier.html>  
(<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.ExtraTreesClassifier.html>)

Gradient Boosting Machine: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html> (<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html>)

AdaBoost: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>  
(<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>)

## Lesson 13

```
In [36]: # Lesson 13

# Save Model Using Pickle
from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
import pickle
url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indian-s-diabetes.data.csv"
names = ['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class']
dataframe = read_csv(url, names=names)
array = dataframe.values
X = array[:,0:8]
Y = array[:,8]
test_size = 0.33
seed = 7
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=test_size,
random_state=seed)
# Fit the model on 67%
model = LogisticRegression(solver='liblinear')
model.fit(X_train, Y_train)
# save the model to disk
filename = 'finalized_model.sav'
pickle.dump(model, open(filename, 'wb'))

# some time later...

# Load the model from disk
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(X_test, Y_test)
print(result)
```

0.7559055118110236

```
In [37]: import pickle

# Save model
# with open('model.pkl', 'wb') as file:
#     pickle.dump(model, file)

# Reload model
# with open('model.pkl', 'rb') as file:
#     model_reloaded = pickle.load(file)
```

In [ ]: